

Real Estate Data Engineering Pipeline – Technical Documentation

1. Overview

This project implements a complete data-engineering workflow for a real-estate company. The goal of the pipeline is to:

- Extract daily **Leads** and **Sales** data
 - Transform the raw data into a Star Schema
 - Load the cleaned and modeled data into a data warehouse (MySQL)
 - Automate all steps using Apache Airflow
-

2. Understanding the Source Data

Before designing the warehouse, I started by exploring the raw Excel sheets:

- **DE LEADS**
- **DE SALES**

I loaded both sheets into temporary MySQL tables using a simple ingestion script.

```
leads_df = pd.read_excel(file_path, sheet_name="DE LEADS")
sales_df = pd.read_excel(file_path, sheet_name="DE SALES")
leads_df.to_sql("de_leads_raw", engine, if_exists="replace", index=False)
sales_df.to_sql("de_sales_raw", engine, if_exists="replace", index=False)
```

This allowed me to understand each column, its datatype, and its business meaning – as much as I could – which helped determine which attributes should become **dimensions** and which metrics should belong to **fact** tables.

3. Environment Setup

Tools & Technologies

- **MySQL** (local)
- **Apache Airflow** (Docker-based)
- **SQLAlchemy**

- Pandas
- OpenPyXL
- MySQL Workbench (for ERD Reverse Engineering)


Challenges

- This was my first time setting up and working with Airflow and I had to configure extra Python packages inside the Airflow container (e.g., pandas, openpyxl, pymysql).
 - Time constraints while designing and testing each step of the pipeline.
-

4. High-Level Architecture

4.1 ERD

I generated the ERD automatically using the **Reverse Engineering** feature in MySQL Workbench after loading the DWH tables.



4.2 Logical flow:

1. **Ingestion (Python / Airflow)**
 - Read Excel → pandas

- Write to `realestate_source.de_leads_raw` & `realestate_source.de_sales_raw`
- 2. **Staging Layer (`realestate_stg`) – `01_staging.sql`**
 - Create:
 - `stg_leads`
 - `stg_sales`
 - These are **cleaned/renamed** versions of raw tables:
 - Rename IDs, basic type/column standardization.
- 3. **DWH Layer (`realestate_dwh`)**
 - **Lead dims – `02_lead_dims.sql`**
 - `dim_lead_status`: funnel mapping, win/loss/active flags.
 - `dim_lead_type`: primary vs resale vs commercial vs referral broker, with `is_commercial` flag.
 - `dim_contact_method`: cleans `method_of_contact` into categories like:
 - `facebook, instagram, search/web, website form, paid ads form, referral, phone/sms, offline, broker/partner`, etc.
 - `dim_lead_source`: normalizes `lead_source` to buckets like `facebook, instagram, google/search, website/form, broker/partner, referral, offline, internal`, etc.
 - `dim_campaign`: cleans raw campaign names, decodes `%20`, classifies into types:
 - `BRAND, REMARKETING, PROSPECTING, SMS, PERFORMANCE_DISPLAY, OTHER.`
 - **Core dims – `03_core_dims.sql`**
 - `dim_date` with:
 - `date_key (YYYYMMDD)`
 - `year, quarter, month, month_name, day_of_month, day_of_week, day_name, week_of_year, is_weekend.`
 - Populated by stored procedure
`populate_dim_date('2010-01-01', '2040-12-31').`
 - `dim_customer` from distinct `customer_id`.
 - `dim_agent` from distinct `user_id`.
 - `dim_area` from `area_id` (leads/sales).
 - `dim_compound` from `compound_id`.
 - `dim_developer` from `developer_id` (for future richer dev analytics).
 - **Sales-specific dims – `04_sales_dims.sql`**
 - `dim_property_type`:
 - Normalizes property type into groups:

- RESIDENTIAL_APT, RESIDENTIAL_VILLA,
VACATION, COMMERCIAL, OTHER.
- Flags: `is_residential`, `is_commercial`, `is_vacation`.
- `dim_sale_category`:
 - Primary, Resale Buyer, Resale Seller, Developer
Resale, Developer Commercial Sale, Navy Now, etc.
 - Groups: PRIMARY, RESALE, DEVELOPER_COMMERCIAL,
INSTANT_BUY, OTHER.
 - Flags: `is_primary`, `is_resale`, `is_developer`,
`is_instant_buy`.

4. Facts

`fact_lead - 05_fact_lead.sql`

- Grain: **1 row per lead_id (latest snapshot)**.
- Surrogate FKS:
 - `customer_key`
 - `agent_key`
 - `lead_status_key`
 - `lead_type_key`
 - `contact_method_key`
 - `lead_source_key`
 - `campaign_key`
 - `area_key`
 - `compound_key`
 - `developer_key`
 - `created_date_key`, `last_request_date_key`,
`last_contact_date_key`
- Measures / attributes:
 - `budget`
 - `is_buyer`, `is_seller`, `is_commercial`
 - `meeting_flag`, `do_not_call_flag`
 - `location`
- Uses `ROW_NUMBER() OVER (PARTITION BY lead_id ORDER BY COALESCE(updated_at, created_at) DESC)` to keep **only the latest version** of each lead.
- Indexes on natural keys in dim tables (`customer_id`, `user_id`, `area_id`, etc.) to speed up joins.

5. `fact_sale - 06_fact_sale.sql`

- Grain: **1 row per sale_id**.
- Links back to:
 - `fact_lead` via `lead_id` (FK).

- `dim_property_type`, `dim_sale_category`, `dim_area`,
`dim_compound`, `dim_date` (reservation/contract).
- Measures / attributes:
 - `unit_value`, `expected_value`, `actual_value`
 - `years_of_payment`
 - `unit_location`

5. Ingestion & Orchestration (Airflow)

- **`ingestion.py`:**
 - Uses `MySqlHook(mysql_conn_id="mysql_source")` to get an SQLAlchemy engine.
 - Reads both Excel sheets via pandas and writes into:
 - `realestate_source.de_leads_raw`
 - `realestate_source.de_sales_raw`
 - Performs simple count checks (source vs DB).
- **`realestate_dag.py`:**
 - Defines DAG `realestate_pipeline`.
 - Uses a helper `run_sql_file(sql_path)` that opens a `.sql` and executes statements one by one with SQLAlchemy.
 - Defines `run_ingestion()` (local version) that also reads Excel and writes into `realestate_source.de_leads_raw` and `de_sales_raw`.
 - Tasks:
 - `ingest_excel` → PythonOperator calling `run_ingestion`
 - `build_staging` → runs `01_staging.sql`
 - `build_lead_dims` → `02_lead_dims.sql`
 - `build_core_dims` → `03_core_dims.sql`
 - `build_sales_dims` → `04_sales_dims.sql`
 - `build_fact_lead` → `05_fact_lead.sql`
 - `build_fact_sale` → `06_fact_sale.sql`
 - Task dependency chain:
`ingest >> stg >> lead_dims >> core_dims >> sales_dims >>`
`fact_lead >> fact_sale`

6. Analytics & KPI Layer

6.1 Lead Funnel Distribution

Shows how leads are distributed across different funnel stages.

```
5    -- Funnel distribution
6 •  SELECT
7      dls.funnel_stage,
8      COUNT(*) AS leads_count
9  FROM fact_lead f
10 JOIN dim_lead_status dls ON dls.lead_status_key = f.lead_status_key
11 GROUP BY dls.funnel_stage
12 ORDER BY leads_count DESC;
13
```

Result Grid	
funnel_stage	leads_count
DISQUALIFIED	15489
CONTACT_ATTEMPT	12764
TO_CONTACT	12484
INTERNAL_REASSIGN	10897
OTHER	1718
MEETING_HELD	1372
WON	604
MEETING_BOOKED	314
RESERVATION	96
CONTRACT	63
MEETING_CANCELED	44

6.2 Funnel Stage Counts

Aggregated counts for each stage of the funnel.


```
14    -- Funnel counts by stage
15 •  SELECT
16      SUM(dls.funnel_stage = 'TO_CONTACT')      AS to_contact,
17      SUM(dls.funnel_stage = 'CONTACT_ATTEMPT') AS contacted,
18      SUM(dls.funnel_stage = 'MEETING_BOOKED')  AS meeting_booked,
19      SUM(dls.funnel_stage = 'MEETING_HELD')    AS meeting_held,
20      SUM(dls.funnel_stage = 'RESERVATION')     AS reservation,
21      SUM(dls.funnel_stage = 'CONTRACT')        AS contract_stage,
22      SUM(dls.funnel_stage = 'WON')              AS sales
23  FROM fact_lead f
24 JOIN dim_lead_status dls ON dls.lead_status_key = f.lead_status_key;
25
```

Result Grid							
to_contact	contacted	meeting_booked	meeting_held	reservation	contract_stage	sales	
12484	12764	314	1372	96	63	604	

6.3 Agent Performance

Summaries leads handled, meetings done, won deals, and win rate.

```
35 • SELECT
36     da.user_id AS agent_id,
37     COUNT(*) AS total_leads,
38     SUM(dls.is_won) AS sales,
39     SUM(dls.funnel_stage = 'MEETING_HELD') AS meetings,
40     ROUND(SUM(dls.is_won) / COUNT(*) * 100, 2) AS win_rate_pct
41 FROM fact_lead f
42 JOIN dim_agent da      ON da.agent_key      = f.agent_key
43 JOIN dim_lead_status dls ON dls.lead_status_key = f.lead_status_key
44 GROUP BY da.user_id
45 ORDER BY sales DESC;
46
```



The screenshot shows a database query result grid titled "Result Grid". The grid has columns: agent_id, total_leads, sales, meetings, and win_rate_pct. The data is sorted by sales in descending order. The first few rows are highlighted in blue.

agent_id	total_leads	sales	meetings	win_rate_pct
2	388	28	2	7.22
150	809	19	2	2.35
239	67	13	0	19.40
135	38	12	1	31.58
222	38	11	6	28.95
442	92	10	8	10.87
189	287	10	8	3.48
336	46	9	0	19.57
181	13	9	0	69.23
387	44	8	2	18.18
259	33	8	10	24.24
156	20	8	1	40.00

6.4 Lead Source Performance

Measures which marketing channels generate the best leads.

```

48 • SELECT
49     dlsrc.source_group,
50     COUNT(*) AS total_leads,
51     SUM(dls.is_won) AS sales,
52     ROUND(SUM(dls.is_won) / COUNT(*) * 100, 2) AS win_rate_pct
53 FROM fact_lead f
54 JOIN dim_lead_source dlsrc ON dlsrc.lead_source_key = f.lead_source_key
55 JOIN dim_lead_status dls  ON dls.lead_status_key    = f.lead_status_key
56 GROUP BY dlsrc.source_group
57 ORDER BY sales DESC;

```

Result Grid | Filter Rows: | Export: | Wrap Cell Content: |

source_group	total_leads	sales	win_rate_pct
OTHER	5482	292	5.33
SOCIAL	32669	117	0.36
SEARCH	9590	89	0.93
OWNED_MEDIA	4499	34	0.76
OFFLINE	2815	29	1.03
REFERRAL	164	27	16.46
PARTNER	155	7	4.52
MESSAGING	166	5	3.01
INTERNAL	15	0	0.00
PORTAL	40	0	0.00

6.5 Sales Financial KPIs

```

110 • SELECT
111     COUNT(*)          AS total_sales,
112     SUM(actual_value) AS total_actual_value,
113     AVG(actual_value) AS avg_actual_value,
114     SUM(expected_value) AS total_expected_value,
115     AVG(expected_value) AS avg_expected_value,
116     SUM(unit_value)    AS total_unit_value,
117     AVG(unit_value)    AS avg_unit_value,
118     AVG(years_of_payment) AS avg_years_of_payment
119 FROM realestate_dwh.fact_sale;
120

```

Result Grid | Filter Rows: | Export: | Wrap Cell Content: |

total_sales	total_actual_value	avg_actual_value	total_expected_value	avg_expected_value	total_unit_value	avg_unit_value	avg_years_of_payment
1567	6405783466.00	8037369.469005	4431103624.00	7653028.711572	9722558203.00	6789495.951816	7.188764

7. Next Enhancements

(1) Make Dimensions & Facts Incremental, Not Full Rebuild

Currently, tables are dropped & recreated.

Instead:

- Use `INSERT ... ON DUPLICATE KEY UPDATE`
- Keep `is_current` flags
- Use incremental logic on `updated_at`

(2) Add Data Quality Checks

Great expectations or SQL-based tests:

- Not null checks
- Unique keys
- Referential integrity
- Value distribution checks